Notes on Adjoint Methods for 18.335

Steven G. Johnson

Created Spring 2006, updated April 30, 2021.

1 Introduction

Given the output $\mathbf{x} \in \mathbb{R}^M$ of a complicated computation, such as the solution of a discretized PDE or some other set of M equations, parameterized by P variables $\mathbf{p}$ (*design parameters*, a.k.a. *control variables* or *decision parameters*), we often wish to compute some function $g(\mathbf{x}, \mathbf{p})$ based on the parameters and the solution. For example, if we solve a wave equation, we might want to know the scattered power in some direction. Or, for a mechanical simulation, we might want to know the load-bearing capacity of the structure. For the output of a neural network, g might be the "loss" function comparing a prediction $\mathbf{x}$ to data. Often, however, we want to know more than just the *value* of g—we also want to know its *gradient* $\frac{dg}{d\mathbf{p}}$. Adjoint methods give an efficient way to evaluate $\frac{dg}{d\mathbf{p}}$, with a cost *independent* of P and usually comparable to the cost of solving for $\mathbf{x}$ *once*. (In the context of neural networks, adjoint methods are known as *backpropagation*; whereas in automatic differentiation they are known as *reverse-mode* differentiation.)

The gradient of g with respect to $\mathbf{p}$ is extremely useful. It gives a measure of the *sensitivity* of our answer to the parameters $\mathbf{p}$ (which may, for example, come from some experimental measurements with some associated uncertainties). Or, we may want to perform an *optimization* of g, picking the $\mathbf{p}$ that produce some desired result; in this case the gradient indicates a useful search direction (e.g. for nonlinear conjugate-gradient optimization). For large-scale optimization problems, the number P of design parameters can be hundreds, thousands, or more—this is common in *shape* or *topology optimization* of PDEs, in which $\mathbf{p}$ controls the placement and shape of arbitrary blobs of different materials constituting a given structure/design [1]. Sometimes, this process is called *inverse design*: finding the problem that yields a given solution instead of the other way around. In *deep learning*, one might have neural networks with millions of degrees of freedom parameterizing artificial "neurons." When $P \gg 1$, the amazing efficiency of adjoint methods makes inverse design and deep learning possible.

When I started these notes in 2006, I hadn't found any textbook description of adjoint methods that I particularly liked, which is part of my motivation for writing up these notes. One introduction can be found in [2], and a more general treatment can be found in [3]. Subsequently, Gil Strang wrote a nice introduction to adjoint methods in his book [4].

2 Linear equations

Suppose that the column-vector $\mathbf{x}$ solves the $M \times M$ linear equation $A\mathbf{x} = \mathbf{b}$ where we take $\mathbf{b}$ and A to be real[1] and to depend in some way on $\mathbf{p}$. To evaluate the gradient directly, we would do

$$\frac{dg}{d\mathbf{p}} = g_{\mathbf{p}} + g_{\mathbf{x}}\mathbf{x}_{\mathbf{p}}$$

where the subscripts indicate partial derivatives ($g_{\mathbf{x}}$ is a row vector, $\mathbf{x}_{\mathbf{p}}$ is an $M \times P$ matrix, etc.). Since g is a given function, $g_{\mathbf{p}}$ and $g_{\mathbf{x}}$ are presumably easy to compute. On the other hand, computing $\mathbf{x}_{\mathbf{p}}$ is hard: evaluating it directly by differentiating $A\mathbf{x} = \mathbf{b}$ by a parameter p_i gives $\mathbf{x}_{p_i} = A^{-1}(\mathbf{b}_{p_i} - A_{p_i}\mathbf{x})$. That is, we would have to solve an $M \times M$ linear equation for P right-hand sides, once for every compont of $\mathbf{p}$; this is impractical if P and M are large.

More explicitly, the problematic term is:

$$g_{\mathbf{x}}\mathbf{x}_{\mathbf{p}} = \underbrace{g_{\mathbf{x}}}_{1 \times M} [\underbrace{A^{-1}}_{M \times M} \underbrace{(\mathbf{b}_{\mathbf{p}} - A_{\mathbf{p}}\mathbf{x})}_{M \times P}] = \underbrace{[g_{\mathbf{x}}A^{-1}]}_{1 \times M} \underbrace{(\mathbf{b}_{\mathbf{p}} - A_{\mathbf{p}}\mathbf{x})}_{M \times P},$$

where $A_{\mathbf{p}}\mathbf{x}$ denotes the $M \times P$ matrix with columns $A_{p_i}\mathbf{x}$ for $i = 1, \dots P$.[2] One way of looking at the difficulty is that in the first equation we multiply a $M \times M$ matrix by a $M \times P$ matrix, which costs $O(M^2P)$ work, or equivalently we have multiplications of A^{-1} by P vectors (i.e., solves of P right-hand sides, which in practice would likely use a

---

[1]This involves no loss of generality, since complex linear equations can always be written as real linear equations of twice the size by taking the real and imaginary parts as separate variables.

[2]Technically, $A_{\mathbf{p}}$ is a rank-3 tensor or "three-dimensional matrix," although it almost certainly isn't stored this way. For example, $A_{p_i}\mathbf{x}$ could be computed for each i separately without saving A_{p_i}. Often, A_{p_i} will be very sparse.

factorization of A or an iterative solver rather than explicitly computing A^{-1}).[3] However, this can be ameliorated simply by *parenthesizing in a different way* [4],[4] as shown in the last expression. If we multiply $\boldsymbol{\lambda}^T = g_{\mathbf{x}} A^{-1}$ first, that corresponds to only a *single* solution of an *adjoint equation*[5]

$$A^T \boldsymbol{\lambda} = g_{\mathbf{x}}^T. \tag{1}$$

and then we multiply a *single* vector $\boldsymbol{\lambda}^T$ by our $M \times P$ matrix for only $\theta(MP)$ work. Putting it all together, we obtain:

$$\left. \frac{dg}{d\mathbf{p}} \right|_{\mathbf{f}=0} = g_{\mathbf{p}} - \boldsymbol{\lambda}^T \mathbf{f}_{\mathbf{p}} = g_{\mathbf{p}} - \boldsymbol{\lambda}^T (A_{\mathbf{p}} \mathbf{x} - \mathbf{b}_{\mathbf{p}}).$$

Again, $A(\mathbf{p})$ and $\mathbf{b}(\mathbf{p})$ are presumably specified analytically and thus $A_{\mathbf{p}}$ and $\mathbf{b}_{\mathbf{p}}$ can easily be computed (in some cases automatically, by automatic program differentiators such as ADIFOR). Note that the adjoint problem is of the same size as the original $A\mathbf{x} = \mathbf{b}$ system, can use the same factorization (e.g. LU factorization $A = LU$ immediately gives $A^T = U^T L^T$), has the same condition number, and has the same spectrum of eigenvalues (the eigenvalues of A and A^T are identical) so iterative algorithms will have similar performance (and can use similar preconditioners)—in every sense, solving the adjoint problem should be no harder than solving the original problem.

3 Nonlinear equations

If $\mathbf{x}$ satisfies some general, possibly nonlinear, equations $\mathbf{f}(\mathbf{x}, \mathbf{p}) = 0$, the process is almost exactly the same. Differentiating the $\mathbf{f}$ equation, we find $\mathbf{f}_{\mathbf{x}} \mathbf{x}_{\mathbf{p}} + \mathbf{f}_{\mathbf{p}} = 0$ and thus $\mathbf{x}_p = -\mathbf{f}_{\mathbf{x}}^{-1} \mathbf{f}_{\mathbf{p}}$. Hence, we write

$$\frac{dg}{d\mathbf{p}} = g_{\mathbf{p}} + g_{\mathbf{x}} \mathbf{x}_{\mathbf{p}} = g_{\mathbf{p}} - \underbrace{g_{\mathbf{x}}}_{1 \times M} \underbrace{[\underbrace{\mathbf{f}_{\mathbf{x}}^{-1}}_{M \times M} \underbrace{\mathbf{f}_{\mathbf{p}}}_{M \times P}]}_{} = g_{\mathbf{p}} - \underbrace{[g_{\mathbf{x}} \mathbf{f}_{\mathbf{x}}^{-1}]}_{1 \times M} \underbrace{\mathbf{f}_{\mathbf{p}}}_{M \times P}.$$

We solve for $\mathbf{x}$ by whatever method, then solve for $\boldsymbol{\lambda}$ from

$$\mathbf{f}_{\mathbf{x}}^T \boldsymbol{\lambda} = g_{\mathbf{x}}^T, \tag{2}$$

and finally obtain

$$\left. \frac{dg}{d\mathbf{p}} \right|_{\mathbf{f}=0} = g_{\mathbf{p}} - \boldsymbol{\lambda}^T \mathbf{f}_{\mathbf{p}}. \tag{3}$$

The only difference is that the adjoint equation (2) is not simply the adjoint of the equation for $\mathbf{x}$. Still, it is a single $M \times M$ linear equation for $\boldsymbol{\lambda}$ that should be of comparable (or lesser) difficulty to solving for $\mathbf{x}$ (being equivalent to the cost of one Newton iteration for $\mathbf{f} = 0$).

4 Eigenproblems

As a more complicated example illustrating the use of equations (2) and (3) from the previous sections, let us suppose that we are solving a linear eigenproblem $A\mathbf{x} = \alpha\mathbf{x}$ and looking at some function $g(\mathbf{x}, \alpha, \mathbf{p})$. For simplicity, assume that A is real-symmetric and that α is simple (non-degenerate; i.e., $\mathbf{x}$ is the only eigenvector for α).[6] In this case, we now have $M + 1$ unknowns described by the column vector:

$$\tilde{\mathbf{x}} = \begin{pmatrix} \mathbf{x} \\ \alpha \end{pmatrix}.$$

The eigenequation $\mathbf{f} = A\mathbf{x} - \alpha\mathbf{x}$ only gives us M equations and doesn't completely determine $\tilde{\mathbf{x}}$, for two reasons. First, of course, there are many possible eigenvalues, but let's assume that we have picked one in some fashion (e.g. the smallest α, or the α closest to π, or the third largest $|\alpha|$, or ...). Second, the eigenequation does not determine the length $|\mathbf{x}|$; let's arbitrarily pick $|\mathbf{x}| = 1$ or $\mathbf{x}^T \mathbf{x} = 1$. This gives us $M + 1$ equations $\tilde{\mathbf{f}} = 0$ where:

$$\tilde{\mathbf{f}} = \begin{pmatrix} \mathbf{f} \\ \mathbf{x}^T \mathbf{x} - 1 \end{pmatrix}.$$

We'll need $M + 1$ adjoint variables $\tilde{\boldsymbol{\lambda}}$:

$$\tilde{\boldsymbol{\lambda}} = \begin{pmatrix} \boldsymbol{\lambda} \\ \beta \end{pmatrix}.$$

The adjoint equations (2) then give:

$$(A - \alpha I)\boldsymbol{\lambda} = g_{\mathbf{x}}^T - 2\beta\mathbf{x}, \tag{4}$$

$$-\mathbf{x}^T \boldsymbol{\lambda} = g_\alpha. \tag{5}$$

The first equation, at first glance, seems to be problematic: $A - \alpha I$ is singular, with a null space of $\mathbf{x}$. It's, okay,

---

[3]If M is sparse, then the cost might be significantly less than this $O(M^2 P)$ upper bound, but in any case solving P right-hand sides will be significantly more costly than solving a single right-hand side for the adjoint formulation.

[4]Another way of looking at this, and the source of the $\boldsymbol{\lambda}$ notation, is to think of sort of a "Lagrange multiplier" process: replace g with $\tilde{g} = g - \boldsymbol{\lambda}^T \mathbf{f}$ by adding a multiple $\boldsymbol{\lambda}$ of $\mathbf{f} = 0$, and then choose $\boldsymbol{\lambda}$ is a clever way to cancel the annoying derivative term. This gives the same result, and may be easier to generalize to some more complicated circumstances, however, such as differential-algebraic equations [3].

[5]For complex-valued $\mathbf{x}$ and A and real g, instead of the transpose A^T one typically obtains the adjoint $A^\dagger = A^{T*}$ (the conjugate-transpose).

[6]Problems involving degenerate eigenvalues occur surprisingly often in optimization of eigenvalues (e.g. when maximizing the minimum eigenvalue of some system), and must be treated with special care. In that case, a generalization of the gradient is required to determine sensitivities or the steepest-descent direction [5], a more elaborate version of what is called *degenerate perturbation theory* in quantum mechanics [**?**].

though! First, we have to choose β so that solutions of equation (4) *exist*: the right-hand side must be orthogonal to $\mathbf{x}$ so that it is not in the null space of $A - \alpha I$. That is, we must have $\mathbf{x}^T(g_{\mathbf{x}}^T - 2\beta \mathbf{x}) = 0$, and thus $\beta = \mathbf{x}^T g_{\mathbf{x}}^T / 2$ (since $\mathbf{x}^T \mathbf{x} = 1$), and therefore $\boldsymbol{\lambda}$ satisfies:

$$(A - \alpha I)\boldsymbol{\lambda} = (1 - \mathbf{x}\mathbf{x}^T)g_{\mathbf{x}}^T = P g_{\mathbf{x}}^T \qquad (6)$$

where $P = 1 - \mathbf{x}\mathbf{x}^T$ is the projection operator into the space orthogonal to $\mathbf{x}$. This equation then has a solution, and in fact it has infinitely many solutions: we can add any multiple of $\mathbf{x}$ to $\boldsymbol{\lambda}$ and still have a solution. Equivalently, we can write $\boldsymbol{\lambda} = \boldsymbol{\lambda}_0 + \gamma \mathbf{x}$ for $\mathbf{x}^T \boldsymbol{\lambda}_0 = 0$ and some γ. Fortunately, γ is determined by (5): $\gamma = -g_{\alpha}$. Finally, with $\boldsymbol{\lambda}_0$ determined by (6),[7] we can find the desired gradient via (3):

$$\left.\frac{dg}{d\mathbf{p}}\right|_{\mathbf{f}=0} = g_{\mathbf{p}} - \boldsymbol{\lambda}^T A_p \mathbf{x} = g_{\mathbf{p}} - \boldsymbol{\lambda}_0^T A_p \mathbf{x} + g_{\alpha} \mathbf{x}^T A_p \mathbf{x}. \quad (7)$$

If we compare with $\frac{dg}{d\mathbf{p}} = g_{\mathbf{p}} + g_{\mathbf{x}}\mathbf{x}_{\mathbf{p}} + g_{\alpha}\alpha_{\mathbf{p}}$, we immediately see that $\alpha_{\mathbf{p}} = \mathbf{x}^T A_p \mathbf{x}$. This is a well-known result from quantum physics and perturbation theory, where it is known as the Hellman-Feynman theorem.

5 Vector–Jacobian products (vJp) and automatic differentiation (AD)

Another way of thinking about adjoint methods is that they correspond to the observation that the **vector–Jacobian product** $\mathbf{v}^T \mathbf{x}_{\mathbf{p}}$ (a "**vJp**"), for *any* given vector $\mathbf{v} \in \mathbb{R}^M$, is much cheaper to compute than the $M \times P$ Jacobian matrix $\mathbf{x}_{\mathbf{p}}$ itself. In our nonlinear-equation setting $\mathbf{f}(\mathbf{x}, \mathbf{p}) = 0$ of Sec. 3, we found that

$$\mathbf{v}^T \mathbf{x}_{\mathbf{p}} = -\boldsymbol{\lambda}^T \mathbf{f}_{\mathbf{p}}$$

where $\boldsymbol{\lambda}$ solves the adjoint equation $\mathbf{f}_{\mathbf{x}}^T \boldsymbol{\lambda} = \mathbf{v}$. Computing the chain rule $g_{\mathbf{x}}\mathbf{x}_{\mathbf{p}}$ ($\mathbf{v} = g_{\mathbf{x}}^T$) is just one application of this general principle.

Nowadays, there are many software tools for "automatic differentiation" or **AD** (such as *autograd* and *JAX* in Python or *Zygote* in Julia): given a computer program to compute $g(p)$, they can *automatically* compute the gradient $\frac{dg}{dp}$ for you, using the analogue of an adjoint method (called "reverse-mode" AD), applying the chain rule from

left-to-right as a sequence of vJp operations. It is extremely useful to have an understanding of adjoint methods even if you are using AD tools, however:

- AD tools **cannot differentiate** many types of code, especially code that calls out to external libraries written in other programming languages.

  - In such cases, however, you can **supply a "manual" vJp** implementation for the computational steps that AD can't handle, and then AD will automatically compose your vJp into the chain rule for all the surrounding pieces.

  - For example, suppose you are performing a composition $g(\mathbf{x}(\mathbf{y}(\mathbf{z}(\mathbf{p}))))$ of several computations involving intermediate vectors $\{\mathbf{x}, \mathbf{y}, \mathbf{z}\}$, and want the derivative $g_{\mathbf{x}}\mathbf{x}_{\mathbf{y}}\mathbf{y}_{\mathbf{z}}\mathbf{z}_{\mathbf{p}}$, but the AD tool doesn't know how to analyze $\mathbf{x}(\mathbf{y})$ because it calls an external Fortran linear-algebra library. You just need to implement a function that computes the vJp $\mathbf{v}^T \mathbf{x}_{\mathbf{y}}$ for any given $\mathbf{v}$ using a "manual" adjoint method, and the AD system will compute the *rest* of the derivatives and put them together with your vJp into a chain rule for $\frac{dg}{d\mathbf{p}}$ automatically.

- AD tools often **generate suboptimal code for iterative/approximate algorithms**.

  - For example, suppose that you have a program solving for $\mathbf{x}$ by Newton's method for a nonlinear problem (or perhaps by a Krylov method, such as GMRES, Arnoldi, etc., for a linear problem). If you apply AD "blindly" to this program, it will **treat it as a generic recurrence** (see my notes on adjoint methods for recurrences) and try to **backpropagate** adjoint variables through the recurrence, which is costly because it requires **storing all intermediate iteration steps** in general. Instead, you can apply the adjoint method (manually, alas) directly to the final *solution*, which only requires one additional linear solve and no knowledge of the intermediate iterations that yielded the solution.

  - Put another way, if you have a function computing $\mathbf{x}(\mathbf{p})$ *approximately* (e.g. by Newton iterations), you often only need the *approximate gradient* as well, as long as you make the error tolerances for both $\mathbf{x}$ and the vJp $\mathbf{v}^T \mathbf{x}_{\mathbf{p}}$ sufficiently small. AD cannot do this, and ends up wasting a lot of effort trying to **exactly differentiate the *error*** in your approximation.

---

[7] Since P commutes with $A - \alpha$, we can solve for $\boldsymbol{\lambda}_0$ easily by an iterative method such as conjugate gradient: if we start with an initial guess orthogonal to $\mathbf{x}$, all subsequent iterates will also be orthogonal to $\mathbf{x}$ and will thus converge to $\boldsymbol{\lambda}_0$ (except for roundoff, which can be corrected by multiplying the final result by P).

6 Initial-value problems

So far, we have looked at $\mathbf{x}$ that are determined by "simple" algebraic equations (which may come from a PDE, etcetera). What if, instead, we are determining $\mathbf{x}$ by integrating a set of equations in *time*? The simplest example of this is an initial-value problem for a linear, time-independent, homogeneous set of ODEs:

$$\dot{\mathbf{x}} = B\mathbf{x}$$

whose solution after a time t for $\mathbf{x}(0) = \mathbf{b}$ is formally:

$$\mathbf{x} = \mathbf{x}(t) = e^{Bt}\mathbf{b}.$$

This, however, is exactly a linear equation $A\mathbf{x} = \mathbf{b}$ with $A = e^{-Bt}$, so we can just quote our results from earlier! That is, suppose we are optimizing (or evaluating the sensitivity) of some function $g(\mathbf{x}, \mathbf{p})$ based on the solution $\mathbf{x}$ at time t. Then we find the adjoint vector $\boldsymbol{\lambda}$ via (1):

$$e^{-B^T t}\boldsymbol{\lambda} = g_{\mathbf{x}}^T.$$

Equivalently, $\boldsymbol{\lambda}$ is the exactly the solution $\boldsymbol{\lambda}(t)$ after a time t of its *own* adjoint ODE:

$$\dot{\boldsymbol{\lambda}} = B^T \boldsymbol{\lambda}$$

with initial condition $\boldsymbol{\lambda}(0) = g_{\mathbf{x}}^T$. We should have expected this by now: solving for $\boldsymbol{\lambda}$ always involves a task of similar complexity to finding $\mathbf{x}$, so if we found $\mathbf{x}$ by integrating an ODE then we find $\boldsymbol{\lambda}$ by an ODE too! Of course, we need not solve these ODEs by matrix exponentials; we can use Runge-Kutta, forward Euler, or (if B comes from a PDE) whatever scheme we deem appropriate (e.g. Crank-Nicolson).

One important property to worry about is *stability*, and here we are in luck. The eigenvalues of B and B^T are the same, and so if one is stable (eigenvalues with real parts ≤ 0) then the other is!

Finally, we can write down the gradient $\frac{dg}{d\mathbf{p}}$ via equation (3):

$$\frac{dg}{d\mathbf{p}} = g_{\mathbf{p}} - \boldsymbol{\lambda}^T (A_{\mathbf{p}}\mathbf{x} - \mathbf{b}_{\mathbf{p}}).$$

Now, since $A = e^{-Bt}$, one might be tempted to write $A_{\mathbf{p}} = -B_{\mathbf{p}}t \cdot A$, but this is not true except in the *very* special case where $B_{\mathbf{p}}$ commutes with B! Unfortunately, the general expression for differentiating a matrix exponential turns out to be more complicated: $A_{\mathbf{p}} = -\int_0^t e^{-Bt'} B_{\mathbf{p}} e^{-B(t-t')} dt'$, and so,

$$\frac{dg}{d\mathbf{p}} = g_{\mathbf{p}} + \int_0^t \boldsymbol{\lambda}^T(t-t') B_{\mathbf{p}}\mathbf{x}(t') dt' + \boldsymbol{\lambda}^T \mathbf{b}_{\mathbf{p}}.$$

This is especially unfortunate because it usually means that we have to *store* $\mathbf{x}(t')$ at all times $0 \leq t' \leq t$ in order to compute the integral. Adjoint methods are storage-intensive for time-dependent problems!

More generally, of course, one might wish to include time-varying A, nonlinearities, inhomogeneous (source) terms, etcetera, into the equations to integrate. A very general formulation of the problem, for differential-algebraic equations (DAEs), can be found in [3]. A similar general principle remains, however: the adjoint variable $\boldsymbol{\lambda}$ is determined by integrating a similar (adjoint) DAE, using the *final* value of $\mathbf{x}(t)$ to compute the *initial* condition of $\boldsymbol{\lambda}(0)$. In fact, the $\boldsymbol{\lambda}(t)$ equation is actually often interpreted as being integrated *backwards* in time from t to 0 (a form of "**backpropagation**"). Alternatively, one can consider a "discrete-time" situation of *recurrence* equations, in which case the adjoint problem is a recurrence "backward in time"—see my online notes on adjoint methods for recurrences.

References

[1] O. Sigmund and K. Maute, "Topology optimization approaches," *Structural and Multidisciplinary Optimization*, vol. 48, pp. 1031–1055, 2013.

[2] R. M. Errico, "What is an adjoint model?," *Bulletin Am. Meteorological Soc.*, vol. 78, pp. 2577–2591, 1997.

[3] Y. Cao, S. Li, L. Petzold, and R. Serban, "Adjoint sensitivity analysis for differential-algebraic equations: The adjoint DAE system and its numerical solution," *SIAM J. Sci. Comput.*, vol. 24, no. 3, pp. 1076–1089, 2003.

[4] G. Strang, *Computational Science and Engineering*. Wellesley, MA: Wellesley-Cambridge Press, 2007.

[5] A. P. Seyranian, E. Lund, and N. Olhoff, "Multiple eigenvalues in structural optimization problems," *Structural Optimization*, vol. 8, pp. 207–227, 1994.

Figure 1: Optimized $V(x)$ (scaled by 1/1000) and $\psi(x)$ for $\psi_0(x) = 1 + \sin[\pi x + \cos(3\pi x)]$ after 500 cg iterations.

7 Example inverse design

As a more concrete example of an inverse-design problem, let's consider the Schrodinger eigen-equation in one dimension,

$$\left[-\frac{d^2}{dx^2} + V(x) \right] \psi(x) = E\psi(x),$$

with periodic boundaries $\psi(x+2) = \psi(x)$. Normally, we take a given $V(x)$ and solve for ψ and E. Now, however, we will specify a particular $\psi_0(x)$ and find the $V(x)$ that gives $\psi(x) \approx \psi_0(x)$ for the ground-state eigenfunction (i.e. for the smallest eigenvalue E). In particular, we will find the $V(x)$ that minimizes

$$g = \int_{-1}^{1} |\psi(x) - \psi_0(x)|^2 dx.$$

To solve this numerically, we will discretize the interval $x \in [-1, 1)$ with M equally-spaced points $x_n = n\Delta x$ ($\Delta x = \frac{2}{M+1}$), and solve for the solution $\psi(x_n)$ at these points, denoted by the vector $\boldsymbol{\psi}$. That is, to compare with the notation of the previous sections, we have the eigenvector $\mathbf{x} = \boldsymbol{\psi}$, the eigenvalue $\alpha = E$, and the parameters $V(x_n)$ or $\mathbf{p} = \mathbf{V}$. If we discretize the eigenoperator with the usual center-difference scheme, we get $A\boldsymbol{\psi} = E\boldsymbol{\psi}$ for:

$$A = \frac{1}{\Delta x^2} \begin{pmatrix} 2 & -1 & 0 & \cdots & 0 & -1 \\ -1 & 2 & -1 & 0 & \cdots & \\ 0 & -1 & 2 & -1 & 0 & \cdots \\ \vdots & & & \ddots & & \\ & & & -1 & 2 & -1 \\ -1 & 0 & \cdots & 0 & -1 & 2 \end{pmatrix} + \mathrm{diag}(\mathbf{V}).$$

As before, we normalize $\boldsymbol{\psi}$ (and $\boldsymbol{\psi}_0$) to $\boldsymbol{\psi}^T \boldsymbol{\psi} = 1$,[8] giving a projection operator $P = 1 - \boldsymbol{\psi}\boldsymbol{\psi}^T$ (or $P = 1 - |\psi\rangle \langle \psi|$, in Dirac notation). The discrete version of g is now

$$g(\boldsymbol{\psi}, \mathbf{V}) = (\boldsymbol{\psi} - \boldsymbol{\psi}_0)^T (\boldsymbol{\psi} - \boldsymbol{\psi}_0)\Delta x$$

where $\boldsymbol{\psi}_0$ is $\psi_0(x_n)$, our target eigenfunction. Therefore, $g_{\boldsymbol{\psi}} = 2(\boldsymbol{\psi} - \boldsymbol{\psi}_0)^T \Delta x$ and thus, by eq. (6), we find $\boldsymbol{\lambda}$ via:

$$(A - E)\boldsymbol{\lambda} = 2P(\boldsymbol{\psi} - \boldsymbol{\psi}_0)\Delta x, \qquad (8)$$

with $P\boldsymbol{\lambda} = 0$ ($\boldsymbol{\lambda} = \boldsymbol{\lambda}_0$ since $g_E = 0$). $g_{\mathbf{V}}$ and g_E are both 0. Moreover, A_{V_n} is simply the matrix with 1 at (n,n) and 0's elsewhere, and thus from (7):

$$\frac{dg}{dV_n} = -\lambda_n \psi_n$$

or equivalently $\frac{dg}{d\mathbf{V}} = -\boldsymbol{\lambda} \odot \boldsymbol{\psi}$ where $\odot$ is the pointwise product (.* in Matlab).

Whew! Now how do we solve these equations numerically? This is illustrated by the Matlab function schrodinger_fd_adj given below. We set up A as a sparse matrix, then find the smallest eigenvalue and eigenvector via the eigs function (which uses an iterative Arnoldi method). Then we solve (8) for $\boldsymbol{\lambda}$ via the Matlab pcg function (preconditioned conjugate-gradient, although we don't bother with a preconditioner).

Then, given g and $\frac{dg}{d\mathbf{V}}$, we then just plug it into some optimization algorithm. In particular, nonlinear conjugate gradient seems to work well for this problem.[9]

7.1 Optimization results

In this section, we give a few example results from running the above procedure (nonlinear cg optimization) for $M = 100$. As the starting guess for our optimization, we'll just use $V(x) = 0$. That is, we are doing a *local optimization* in a *100-dimensional space*, using the adjoint method to get the gradient. It is somewhat remarkable that this works—in a few seconds on a PC, it converges to a very good solution!

We'll try a couple of example $\psi_0(x)$ functions. To start with, let's do $\psi_0(x) = 1 + \sin[\pi x + \cos(3\pi x)]$. (Note that the ground-state ψ will never have any nodes, so we require $\psi_0 \geq 0$ everywhere.) This $\psi_0(x)$, along with the resulting $\psi(x)$ and $V(x)$ after 500 cg iterations, are shown in figure 1. The solution $\psi(x)$ matches $\psi_0(x)$ very well except for a couple of small ripples, and $V(x)$ is quite complicated—not something you could easily guess!

---

[8]We also have an arbitrary choice of sign, which we fix by choosing $\int \psi dx > 0$.

[9]I used the nonlinear conjugate-gradient Matlab conj_grad routine from:

    http://www2.imm.dtu.dk/~hbn/Software/

Oh, but that ψ_0 was too easy! Let's try one with discontinuities: $\psi_0(x) = 1 - |x|$ for $|x| < 0.5$ and 0 otherwise (which looks a bit like a "house"). This $\psi_0(x)$, along with the resulting $\psi(x)$ and $V(x)$ after 500 cg iterations, are shown in figure 2. Amazingly, it still captures ψ_0 pretty well, although it has a bit more trouble with the discontinuities than with the slope discontinuity. This time, we let it converge for 5000 cg iterations to give it a bit more time. Was this really necessary? In figure 3, we plot $\psi(x)$ for 10, 20, 40, 80, 160, 320, and 5000 cg iterations. It gets the rough shape pretty quickly, but the discontinuous features are converging fairly slowly. (Presumably this could be improved if we found a good preconditioner, or perhaps by a different optimization method or objective function.)

Figure 2: Optimized $V(x)$ (scaled by 1/10000) and $\psi(x)$ for $\psi_0(x) = 1 - |x|$ for $|x| < 0.5$, after 5000 cg iterations.


Figure 3: Optimized $\psi(x)$ for $\psi_0(x) = 1 - |x|$ for $|x| < 0.5$, after various numbers of nonlinear conjugate-gradient iterations (from 10 to 10000).

6

7.2 Matlab code

The following code solves for g and $\frac{dg}{d\mathbf{V}}$, not to mention
the eigenfunction $\boldsymbol{\psi}$ and the corresponding eigenvalue E,
for a given $\mathbf{V}$ and $\boldsymbol{\psi}_0$.

```
% Usage: [g,gp,E,psi] = schrodinger_fd_adj(x, V, psi0)
%
% Given a column-vector x(:) of N equally spaced x points and a column-vector
% V of the potential V(x) at those points, solves Schrodinger's eigen-equation
%                  [ -d^2/dx^2 + V(x) ] psi(x) = E psi(x)
% with periodic boundaries for the lowest "ground state" eigenvalue E and
% wavefunction psi.
%
% Furthermore, it computes the function g = integral |psi - psi0|^2 and
% the gradient gp = dg/dV (at each point x).

function [g,gp,E,psi] = schrodinger_fd_adj(x, V, psi0)
  dx = x(2) - x(1);
  N = length(x);
  A = spdiags([ones(N,1), -2 * ones(N,1), ones(N,1)], -1:1, N,N);
  A(1,N) = 1;
  A(N,1) = 1;
  A = - A / dx^2 + spdiags(V, 0, N,N);

  opts.disp = 0;
  [psi,E] = eigs(A, 1, 'sa', opts);
  E = E(1,1);
  if sum(psi) < 0
    psi = -psi; % pick sign; note that psi' * psi = 1 from eigs already
  end

  gpsi = psi - psi0;
  g = gpsi' * gpsi * dx;
  gpsi = gpsi * 2*dx;

  P = @(x) x - psi * (psi' * x); % projection onto direction normal to psi

  [lambda,flag] = pcg(A - spdiags(E*ones(N,1), 0, N,N), P(gpsi), 1e-6, N);
  lambda = P(lambda);
  gp = -real(conj(lambda) .* psi);

  disp(g);
```